

Intake triage

AI reads the enquiry. Company rules choose the recipient.

THE CASE

Each week, a professional-services firm receives 40-60 enquiries. A junior analyst spends about eight hours reading them, identifying the work, estimating complexity and deciding who should receive it. This prototype uses AI to identify the work and its complexity, judge whether it is safe to route, and pass it to company workflow rules, which select a person or human-review queue. We tested it through Meridian Advisory, a realistic fictional generalist consultancy, using 30 synthetic enquiries covering a broad range of situations. The AI understood the requested work well. Its main weakness was recognizing when to avoid automatic routing.

1 The firm


Meridian Advisory has six practices, each with a practice lead and a senior lead. Simple organizational rules decide who receives the work.

Strategy & Transformation
Operations & Process
Data & AI

Technology & Systems
Risk & Compliance
Finance & Transactions

2 How the system works

The AI interprets the enquiry; Meridian's workflow rules decide who receives it.



S07, preserved first run. Simple work can still go to a senior lead because the client is an enterprise account.

3 How we tested it

30 fictional enquiries across varied industries. Eighteen routine cases cover simple, moderate and complex work in all six practices; six harder cases mix practice vocabulary but still have a clear owner. Three are ambiguous, two lack or contradict key facts, and one requests unsupported work.

We first defined each client scenario and intended answer. A separate AI model, **not shown the answer**, turned the scenario into realistic client prose. We reviewed it for fidelity, then tested the classifier. Its first results are preserved without tuning against the failures.

Deliberate traps: simple work at a large firm; complex projects at a small one; AI-enabled Operations work; simple regulated work; cross-practice ownership; conflicting form and free-text details.

27 / 30

Automatically routed

24 / 27

Automatic routes correct

25 / 25

Primary service lines correct*

3 / 6

Review cases caught

The first run understood the work well, but missed half the cases needing review. This deliberately difficult synthetic test does not estimate production reliability. *Includes only cases with an intended primary service line.

4 A revealing failure

I02 / OBSERVED FIRST RUN

A client asked for a CRM migration in one business unit. The work was clear; the company size was not.

FORM FIELD / COMPANY SIZE	DESCRIPTION / EXCERPT
1-99 employees	“...the company has about 15,000 employees across twelve countries... ”

AI UNDERSTOOD

- ✓ Technology & Systems
- ✓ Moderate complexity

AI MISSED

- ✗ Conflicting company-size information
- ✗ Marked the case safe to route

ACTUAL ROUTE

Alex Romero

SHOULD HAVE GONE TO

Central Intake Review

Company size changes the recipient: small + moderate goes to the practice lead; enterprise + moderate goes to the senior lead. With those facts in conflict, a human should resolve the contradiction before anyone is assigned.

5 Economics & production rollout


ASSUMPTIONS / COST SENSITIVITY

Eight hours / 50 midpoint enquiries implies **9.6 minutes of manual triage per enquiry**. Assume each review costs one manual triage; a wrong route may cost several times that downstream.

MISROUTE COST	SYSTEM COST VS. MANUAL
1×	20%
3×	40%
10×	110%

With three reviews and three wrong routes, break-even is about **9× one manual triage**. This is sensitivity analysis, not a production ROI forecast. Real economics need actual case mix, review time, misroute cost, response/conversion effects and model cost, excluded here.

PROPOSAL / START IN SHADOW MODE



Track overrides, later reassignment, review rate, conflicting metadata and response time / outcomes.

The analyst keeps routing authority while AI runs beside the existing process. Enable automation only for kinds of cases whose observed real-world error rate is acceptable. Operational/API failures and uncertain cases return to the existing human intake queue.

6 What we learned / next step

The AI understands incoming work well, but needs more evidence before routing every enquiry unattended. Next, learn from real enquiries and analyst corrections, clarify practice ownership, and improve when humans are asked to review. **Expand automation only where observed behavior shows it is safe and valuable.**

INSIDE THE ZIP

Runnable prototype + web demo · firm, practice and routing configuration · 30 synthetic enquiries
Preserved first-run evidence · README · EVALUATION.md · SPEC.md